

# DB-Independent Dynamic CSV Ingestion Framework

**Audience:** Backend / Platform Engineers

**Goal:** Build a production-grade, DB-agnostic ingestion framework that can handle **unknown CSV schemas**, **schema drift**, **multiple RDBMS**, and **future Delta Lake writes** using **Spring Batch**.

## 1. Problem Statement

Business uploads CSV files with **unknown and changing schemas**: - Today: `master_data.csv` with 10 columns - Tomorrow: same file with 3 new columns - Later: data type change (STRING → INT)

Target systems: - MySQL - PostgreSQL - Oracle - SQL Server - (Future) Delta Lake


Key constraints: - Restartable ingestion - Safe schema evolution - No vendor lock-in - Production observability


## 2. High-Level Solution

We use **Spring Batch as an orchestrator**, not as a DB abstraction.

Core ideas: - Canonical schema (DB-neutral) - Dialect-based DDL/DML generation - Pluggable writers (RDBMS / Delta) - Step-wise execution (DDL before DML)

## 3. High-Level Architecture





## 4. Canonical Schema (Foundation)

All schemas are converted into a **canonical representation**.

### 4.1 Canonical Types

```
STRING | INTEGER | LONG | DECIMAL | BOOLEAN | DATE | TIMESTAMP
```

### 4.2 Canonical Column Model

```
class CanonicalColumn {
    String name;
    CanonicalType type;
    Integer length;
    Integer precision;
    Integer scale;
    boolean nullable;
}
```

Why this matters:

- Shields business logic from DB specifics
- Makes Delta Lake mapping trivial
- Enables safe schema comparison

## 5. Step 1 – Schema Discovery (CSV)


### Purpose

- Discover column names
- Infer data types
- Detect schema drift

## How it works

1. Read CSV header
2. Sample first N rows
3. Infer canonical types
4. Store schema in `ExecutionContext`

## Flow



## Notes

- This step is DB-independent
- No table access here
- Safe to re-run

## 6. Step 2 – Schema Management (DDL)

### Purpose


- Create table if missing
- Add new columns
- Modify existing columns (controlled)

### Dialect Pattern

```
interface DatabaseDialect {  
    String createTableSql(String table, List<CanonicalColumn> cols);  
    List<String> alterTableSql(  
        String table,  
        List<CanonicalColumn> csvSchema,  
        List<CanonicalColumn> dbSchema  
    );  
    String upsertSql(String table, Set<String> columns);  
}
```

Each DB has its own implementation: - MySQLDialect - PostgresDialect - OracleDialect - SqlServerDialect

## Flow



## Safety Rules (Recommended)

- Never auto-drop columns
- No narrowing conversions (STRING → INT)
- Allow widening (INT → STRING)
- Feature-flag ALTER operations

## 7. Step 3 – Data Load (Insert / Update)


### Challenges


- Unknown columns
- DB-specific UPSERT syntax
- Batch performance

### Solution

- Read rows as `Map<String, Object>`
- Generate SQL dynamically using dialect
- Batch using JDBC or Spark (Delta)

## Flow





## 8. Writer Abstraction

### Unified Contract

```
interface DataWriter<T> {  
    void write(List<T> items);  
}
```

### RDBMS Writer

- Uses JDBC batch
- Uses dialect UPSERT SQL
- Fully transactional

### Delta Lake Writer

- Uses Spark
- `mergeSchema = true`
- Append / merge supported

This allows **writer swap without job changes**.

## 9. Job Configuration Overview


```
Job: dynamicCsvIngestionJob  
  
Step 1: schemaDiscoveryStep  
Step 2: schemaSyncStep  
Step 3: dataLoadStep
```

Each step: - Commits independently - Restartable - Auditable

## 10. Restartability & Failure Handling

Spring Batch automatically stores: - Last processed row - Step execution status - Failure reason

### Restart Scenario



No manual tracking required.

---

## 11. Observability & Auditing

Recommended tables: - `batch_job_execution` - `batch_step_execution` - `ingestion_file_audit`

Track: - File name - Schema hash - Row count - Start / end time - Target system

---

## 12. Why Spring Batch (Final Justification)

Capability	JDBC Batch	Spring Batch
Dynamic schema	⚠ Manual	✓ Structured
Restartability	✗	✓
Multi-DB	✗	✓
DDL orchestration	✗	✓
Delta Lake ready	✗	✓
Observability	✗	✓

---

## 13. Recommended Next Enhancements

- Schema versioning table

- Staging tables + merge
  - File partitioning (parallelism)
  - Prometheus metrics
  - Iceberg / Hudi writers
- 

## 14. Summary

This framework:

- Handles schema drift safely
- Is DB-agnostic by design
- Scales from RDBMS to Lakehouse
- Is production-ready and extensible

**Spring Batch = Orchestrator**

**Dialect = Portability**

**Canonical Schema = Stability**

---

*End of document*